

# CAB320 - 2023 Exam Prep Guide

Jaden Ussher

July 25, 2024

## Question 1: Edit Distance

The edit distance problem involves finding the minimum number of operations required to transform one string into another. The allowed operations are insertions, deletions, and substitutions (mismatches).

### Steps to Solve Edit Distance using Dynamic Programming

1. **Define the DP Table:** Let  $dp[i][j]$  represent the minimum edit distance between the first  $i$  characters of string  $s1$  and the first  $j$  characters of string  $s2$ .

2. **Initialize the DP Table:** -  $dp[0][0] = 0$  -  $dp[i][0] = i$  for all  $i$  (cost of deleting all characters from  $s1$ ) -  $dp[0][j] = j$  for all  $j$  (cost of inserting all characters into  $s1$ )

$$dp[i][0] = i \quad \text{for } 0 \leq i \leq m$$

$$dp[0][j] = j \quad \text{for } 0 \leq j \leq n$$

3. **Fill the DP Table:** For each pair of indices  $(i, j)$ : - If  $s1[i-1] == s2[j-1]$ , then  $dp[i][j] = dp[i-1][j-1]$  (no operation needed) - Else, choose the minimum cost among:

$$dp[i][j] = \min \begin{cases} dp[i-1][j] + 1 & \text{(deletion)} \\ dp[i][j-1] + 1 & \text{(insertion)} \\ dp[i-1][j-1] + 1 & \text{(substitution)} \end{cases}$$

for i from 1 to m:

  for j from 1 to n:

    if  $s1[i-1] == s2[j-1]$ :

$dp[i][j] = dp[i-1][j-1]$

    else:

$dp[i][j] = \min(dp[i-1][j] + 1, dp[i][j-1] + 1, dp[i-1][j-1] + 1)$

4. **Extract the Result:** The value  $dp[m][n]$  contains the minimum edit distance between  $s1$  and  $s2$ .

### Example Problem

Calculate the edit distance between the strings "kitten" and "sitting".

*Solution*

1. **Initialization:**

	$\emptyset$	s	i	t	t	i	n	g
$\emptyset$	0	1	2	3	4	5	6	7
k	1	0	0	0	0	0	0	0
i	2	0	0	0	0	0	0	0
t	3	0	0	0	0	0	0	0
t	4	0	0	0	0	0	0	0
e	5	0	0	0	0	0	0	0
n	6	0	0	0	0	0	0	0

2. **Filling the DP Table:**

	$\emptyset$	s	i	t	t	i	n	g
$\emptyset$	0	1	2	3	4	5	6	7
k	1	1	2	3	4	5	6	7
i	2	2	1	2	3	4	5	6
t	3	3	2	1	2	3	4	5
t	4	4	3	2	1	2	3	4
e	5	5	4	3	2	2	3	4
n	6	6	5	4	3	3	2	3

3. **Result:** The minimum edit distance is  $dp[6][7] = 3$ .

## Question 2: Search Problems

### Defining State Representation

When defining a state representation for a problem: - Identify all the relevant attributes that uniquely determine the state. - Ensure the state representation is complete and minimal.

**Example:** For a puzzle game, a state might include: - The position of each piece on the board. - The number of moves taken so far.

### Example Problem

Define a state representation for the 8-puzzle problem.

*Solution*

A state can be represented by: - A 3x3 matrix (or a 1D array of length 9) representing the position of each tile. - The position of the blank tile.

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & \emptyset \end{bmatrix}$$

State representation: - Array: [1, 2, 3, 4, 5, 6, 7, 8, 0] - Blank tile position: 8 (indexing from 0)

### Admissible Heuristics

An admissible heuristic  $h(s)$  is one that never overestimates the cost to reach the goal from state  $s$ .

**Example:** For the 8-puzzle problem, the Manhattan distance is an admissible heuristic.

### Example Problem

Prove that the Manhattan distance is an admissible heuristic for the 8-puzzle problem.

*Solution*

The Manhattan distance is the sum of the absolute values of the differences in the horizontal and vertical positions of the tiles from their goal positions.

For each tile  $i$ :

$$h_i(s) = |x_i - x_i^*| + |y_i - y_i^*|$$

where  $(x_i, y_i)$  is the current position of tile  $i$  and  $(x_i^*, y_i^*)$  is the goal position of tile  $i$ .

Since each move of a tile to its goal position contributes exactly 1 to the Manhattan distance and no move can decrease the distance by more than 1, the Manhattan distance never overestimates the cost.

## Question 3: Reinforcement Learning

### Tabular Methods

Understand the Bellman equations and how they apply to value functions and Q-functions.

**Value Function:**

$$V(s) = \max_a Q(s, a)$$

**Q-Learning Update Rule:**

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[ r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$$

**SARSA Update Rule:**

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]$$

### Example Problem

Consider a simple MDP with three states: A, B, C. The transition and reward structure is as follows: - From A, you can go to B with reward -1. - From B, you can go to C with reward 2. - From C, you can go to A with reward 0. - From any state, you can stay in the same state with reward -1.

Apply one step of the Q-learning update rule with  $\alpha = 0.1$  and  $\gamma = 0.9$  for the transition from state A to state B.

*Solution*

Assume current Q-values are:

$$Q(A, B) = 0.5, \quad Q(B, C) = 0.6, \quad Q(C, A) = 0.4$$

Step-by-step: 1. Observe the transition from A to B with reward  $r = -1$ . 2. Calculate the updated Q-value for  $Q(A, B)$ :

$$Q(A, B) \leftarrow 0.5 + 0.1 \left[ -1 + 0.9 \cdot \max_{a'} Q(B, a') - 0.5 \right]$$

Assume  $\max_{a'} Q(B, a') = Q(B, C) = 0.6$ :

$$Q(A, B) \leftarrow 0.5 + 0.1 [-1 + 0.9 \cdot 0.6 - 0.5]$$

$$Q(A, B) \leftarrow 0.5 + 0.1 [-1 + 0.54 - 0.5]$$

$$Q(A, B) \leftarrow 0.5 + 0.1 [-0.96]$$

$$Q(A, B) \leftarrow 0.5 - 0.096$$

$$Q(A, B) \leftarrow 0.404$$

## Question 4: Bayes Rule

Bayes Rule is used to update the probability estimate for a hypothesis given new evidence.

$$P(H|E) = \frac{P(E|H) \cdot P(H)}{P(E)}$$

### Example Problem

A medical test for a disease is 99% sensitive (true positive rate) and 95% specific (true negative rate). The disease prevalence is 1

*Solution*

1. Define the events: -  $D$ : Having the disease -  $T$ : Positive test result
2. Given: -  $P(T|D) = 0.99$  -  $P(T|\neg D) = 1 - 0.95 = 0.05$  -  $P(D) = 0.01$  -  $P(\neg D) = 0.99$
3. Use Bayes' Rule:

$$P(D|T) = \frac{P(T|D) \cdot P(D)}{P(T)}$$

4. Calculate  $P(T)$ :

$$P(T) = P(T|D) \cdot P(D) + P(T|\neg D) \cdot P(\neg D)$$

$$P(T) = 0.99 \cdot 0.01 + 0.05 \cdot 0.99$$

$$P(T) = 0.0099 + 0.0495 = 0.0594$$

5. Calculate  $P(D|T)$ :

$$P(D|T) = \frac{0.99 \cdot 0.01}{0.0594}$$

$$P(D|T) = \frac{0.0099}{0.0594}$$

$$P(D|T) \approx 0.1667$$

So, the probability that a person has the disease given a positive test result is approximately 16.67

## Q-Learning Example

### Problem Description:

Given the following map where the player starts on tile 2 and ends on tile 6. The player receives 10 points for reaching tile 6 and -1 points for all other tiles. From each tile, the player can choose to go (up, down, left, right), but cannot move outside the given tiles. The  $Q(s, a)$  values have been provided:

- $Q(1, \text{down}) = 4$
- $Q(1, \text{right}) = 6$
- $Q(2, \text{left}) = 2$
- $Q(2, \text{right}) = 4$
- $Q(2, \text{down}) = 3$
- $Q(3, \text{left}) = 5$
- $Q(3, \text{down}) = 7$
- $Q(4, \text{right}) = 4$
- $Q(4, \text{up}) = 8$
- $Q(5, \text{left}) = 5$
- $Q(5, \text{right}) = 6$
- $Q(5, \text{up}) = 8$

### Grid Layout:

1	2	3
4	5	6

### Task:

1. Set up a Q-table with the corresponding values.
2. With learning rate  $\alpha = 0.25$  and discount rate  $\gamma = 0.75$ , if the player travels down first then right, what are the updated values in the Q-table?

## Solution

### 1. Q-Table Setup

$Q =$

State	Action	Q-Value
1	down	4
1	right	6
2	left	2
2	right	4
2	down	3
3	left	5
3	down	7
4	right	4
4	up	8
5	left	5
5	right	6
5	up	8

### 2. Q-Learning Update Process

#### Definitions:

- Learning rate  $\alpha = 0.25$
- Discount rate  $\gamma = 0.75$

**Observations:**

1. Initial State: State 2
2. Action: Down (to State 5)
3. Reward: -1 (since it's not tile 6)
4. Next State: State 5
5. Next Action: Right (to State 6)
6. Next Reward: +10 (for reaching tile 6)

*Step 1: From State 2, Action Down*

$$Q(2, \text{down}) \leftarrow Q(2, \text{down}) + \alpha \left[ \text{Reward} + \gamma \max_{a'} Q(5, a') - Q(2, \text{down}) \right]$$

Where:

- Reward = -1
- $\max_{a'} Q(5, a') = \max(5, 6, 8) = 8$  (from the Q-values of state 5)

$$Q(2, \text{down}) \leftarrow 3 + 0.25 [-1 + 0.75 \cdot 8 - 3]$$

$$Q(2, \text{down}) \leftarrow 3 + 0.25 [-1 + 6 - 3]$$

$$Q(2, \text{down}) \leftarrow 3 + 0.25 [2]$$

$$Q(2, \text{down}) \leftarrow 3 + 0.5 = 3.5$$

*Step 2: From State 5, Action Right*

$$Q(5, \text{right}) \leftarrow Q(5, \text{right}) + \alpha \left[ \text{Reward} + \gamma \max_{a'} Q(6, a') - Q(5, \text{right}) \right]$$

Where:

- Reward = 10
- $\max_{a'} Q(6, a') = 0$  (assuming Q-values for state 6 are initially 0 since it's the goal)

$$Q(5, \text{right}) \leftarrow 6 + 0.25 [10 + 0.75 \cdot 0 - 6]$$

$$Q(5, \text{right}) \leftarrow 6 + 0.25 [10 - 6]$$

$$Q(5, \text{right}) \leftarrow 6 + 0.25 [4]$$

$$Q(5, \text{right}) \leftarrow 6 + 1 = 7$$

**Final Updated Q-Table**

State	Action	Q-Value
1	down	4
1	right	6
2	left	2
2	right	4
2	down	3.5
3	left	5
3	down	7
4	right	4
4	up	8
5	left	5
5	right	7
5	up	8

## Summary

The Q-Learning algorithm updates the Q-values based on the observed transitions and rewards. By iterating through multiple episodes and updating the Q-table, the agent learns the optimal policy to maximize the cumulative reward.